

CoE202

Fundamentals of Artificial intelligence

<Big Data Analysis and Machine Learning>

Generalization

Prof. Young-Gyu Yoon
School of EE, KAIST

Contents


- Model selection problem
- Bias-variance tradeoff
- Generalization & Validation
- Splitting dataset
- Model capacity

Recap: Linear classification

For a data set $\mathcal{D} = \{(\vec{x}_1, y_1), (\vec{x}_2, y_2), \dots, (\vec{x}_N, y_N)\}$

Seeks a function $f(\vec{x}; \theta_0, \vec{\theta}) = \vec{\theta} \cdot \vec{x} + \theta_0$

Such that a loss function
 $\mathcal{L} = -\frac{1}{N} \sum_{i=1}^N (y_i \log(f(x_i)) + (1 - y_i) \log(1 - f(x_i)))$
is minimized



Recap: Softmax function

Standard softmax


$$\sigma(z_i) = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$$

Softmax with a smoothness parameter

$$\sigma(z_i) = \frac{e^{\beta z_i}}{\sum_{j=1}^K e^{\beta z_j}}$$

Vector representation of above

$$\sigma(\vec{z}) = \frac{e^{\beta \vec{z}}}{\sum_{j=1}^K e^{\beta z_j}}$$




- This softmax function allows us to convert a meaningless set of scores to a set of probabilities (which is a lot more interpretable!)

Recap: sigmoid function

Sigmoid function

$$S(z) = \frac{e^z}{e^z + 1} = \frac{1}{1 + e^{-z}}$$



- Sigmoid function
 - is monotonic
 - maps all real values to $(0,1)$ → can be used to represent a probability

Recap: Binary cross entropy loss

- **Binary cross entropy (BCE) loss**

$$\mathcal{L} = -(y \log(f(x)) + (1 - y) \log(1 - f(x)))$$

- measures the similarity of two probabilities
- Consider two cases
 - When $y_i=0$, then $f(x_i)=0$ & $y_i=1$, then $f(x_i)=1$
 - When $y_i=0$, then $f(x_i)=1$ & $y_i=1$, then $f(x_i)=0$

- **BCE loss of multiple data samples**

$$\mathcal{L} = -\frac{1}{N} \sum_{i=1}^N (y_i \log(f(x_i)) + (1 - y_i) \log(1 - f(x_i)))$$

Recap: Cross-entropy loss

- Cross entropy (CE) loss for K-class classification

$$\mathcal{L} = - \sum_{j=1}^K y_j \log(f(x)_j)$$

Why is this here?


- When we have multiple data samples, we take an average

$$\mathcal{L} = - \frac{1}{N} \sum_{i=1}^N \sum_{j=1}^K y_{ij} \log(f(x_i)_j)$$

sample index





class index

Recap: Gradient Descent for linear classification




Two trials with different initial parameters

Recap: Gradient Descent for linear classification




Two trials with different initial parameters

Model selection problem




$$f(x) = 0.6934x + 0.6747$$




$$f(x) = -0.0805x^2 + 1.3331x - 0.1339$$

Which one is better?

Model selection problem




$$f(x) = 0.6934x + 0.6747$$




$$f(x) = -0.075x^3 + 0.85x^2 - 1.925x + 2.65$$

Which one is better?


Model selection problem



$$f(x) = 0.6934x + 0.6747$$



$$f(x) = -0.0805x^2 + 1.3331x - 0.1339$$



$$f(x) = -0.075x^3 + 0.85x^2 - 1.925x + 2.65$$

- For a same given data set...
- Higher order polynomial has more degree of freedom
 - 2nd order polynomial can be considered a just special case of 3rd order polynomial
- Higher order polynomial has lower loss value
 - Does it mean it is better?

What is the purpose of machine learning?

- In case of supervised learning, we want our algorithm to learn a function that can represent the relation between the input and the output ... and **we are expecting the function to be useful for unseen data**
- In other words, we are assuming that our data has some **underlying structure** (although it may not be apparent) and we are **trying find the structure** (via function approximation)
- Minimizing the loss function is just a mean, but not our real goal

What does this mean?

- If we know something about the underlying structure of the data, we can (and should) exploit that knowledge
- Choosing an appropriate loss function is important
- We need a measure of our algorithm's **performance for unseen data**

Bias-Variance Tradeoff

- Let's assume that our data is generated as follow

$$y = f(x) + \epsilon$$

- $f(x)$: true relation between the input and the output
- ϵ : noise with zero mean and variance of σ^2
- Here, we want to find $\hat{f}(x)$ that approximates the true function $f(x)$ by minimizing the MSE loss

For a data set $\mathcal{D} = \{(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)\}$

Seeks a function $f : X \rightarrow Y$

Such that a loss function $\mathcal{L} : X \times Y \rightarrow \mathcal{R}$ is minimized

Bias-Variance Tradeoff

- For which ever function $\hat{f}(x)$ we use, the expected error on an unseen sample x is as follow

$$E[(y - \hat{f}(x))^2] = (Bias[\hat{f}(x)])^2 + Var[\hat{f}(x)] + \sigma^2$$

$$Bias[\hat{f}(x)] = E[\hat{f}(x) - y]$$

: “consistent” error (consistently over-estimate or under-estimate)

$$Var[\hat{f}(x)] = E[\hat{f}(x)^2] - E[\hat{f}(x)]^2$$

: fluctuation of $\hat{f}(x)$

$$\sigma^2$$

: irreducible error (even when $\hat{f}(x) = f(x)$)

Bias-Variance Tradeoff

- Bias

$$Bias[\hat{f}(x)] = E[\hat{f}(x) - y]$$



- The bias error is an error from erroneous assumptions in the learning algorithm. High bias can cause an algorithm to **miss the relevant relations** between features and target outputs (underfitting).

- Variance


$$Var[\hat{f}(x)] = E[\hat{f}(x)^2] - E[\hat{f}(x)]^2$$

- The variance is an error from sensitivity to small fluctuations in the training set. High variance can cause an algorithm to **model the random noise** in the training data, rather than the intended outputs (overfitting)

Bias-Variance Tradeoff




Bias-Variance Tradeoff




Overfitting

In classification



In regression




Generalization

- **Generalization** refers to a model's ability to adapt properly to new, previously **unseen data**, drawn from the same distribution as the one used to create the model.
- If performance on training data \approx performance on unseen data
 - then, we can speculate that the model has learned “real” underlying structure of the data set
 - in such case, we say the model generalizes well
- If performance on training data \gg performance on unseen data
 - then, we can speculate that the model has learned something meaningless
 - in such case, we say the model generalizes poorly (or we say the model **overfits** the data)

How do we know if our model “will” generalize well?

- Q) does figure below truly shows overfitting?



For good generalization ...

- From earlier slides, we learned that ...
 - Bias
 - The bias error is an error from erroneous assumptions in the learning algorithm. High bias can cause an algorithm to **miss the relevant relations** between features and target outputs (underfitting).
 - Variance
 - The variance is an error from sensitivity to small fluctuations in the training set. High variance can cause an algorithm to **model the random noise** in the training data, rather than the intended outputs (overfitting)
- **This tells us we have to balance between bias and variance**
- **to not miss the relevant relation**
- **to not model the random noise**
- **But, how do we know if it's a relevant relation or noise?**

How do we know if our model “will” generalize well?

- **Short answer)** Test our performance on new data set!
 - What we really want to know is our model’s prediction performance on unseen new data set...which can be estimated by testing its performance on unseen new data set
 - But then do we have to get truly new data sets?
- Trick: we split our available data into multiple sets
 - We can regard some of them as “old” and some as “new”
 - We train our model with a part of the data set, and then **validate** our model with the other (exclusive) part of the data set

Validation


- We need a method (or measure) to check how well our model generalizes
- An unbiased evaluation of a model fit on the dataset while tuning model hyperparameters
 - The evaluation becomes biased as we tune hyperparameters for better validation

Training, Validation, Test

- How can we check the **performance for unseen data**?
- We (a) split our data, then (b) use only part of it for training and (c) see its performance on the other part of the data
- **How do we split?**
 - **Training data set**: data set that is directly used for parameters updates (i.e., for gradient calculation)
 - **Validation data set**: data set that is indirectly used for training (i.e., for tuning ***hyperparameters**)
 - **Test data set**: data set that is used for final evaluation of the performance. Should never be used anyhow for training

*hyperparameters: parameters whose values are used to control the learning process
(e.g., learning rate, mini-batch size, number of iterations, model complexity, etc)

Training, Validation, Test



- **Depending on the availability and characteristics of the data set, the data can be split differently**
 - **80-10-10**
 - **50-25-25**
 - **Other ratios ...**

Training, Validation, Test



Training data set



Validation data set




Test data set

- (Randomly) split the entire data set into training, validation and test data sets
- Use training set (for parameter updates) and validation set (for hyperparameter tuning) for training
- Use test set for final performance evaluation

Training & Validation

- Typical learning curve



Validation vs. Test

- **How is a test different from validation?**

- From the earlier slide ...“The evaluation becomes biased as we tune hyperparameters for better validation”
- Although validation data set was not directly used for parameter updates, we (human) reflected the results from validation set to tune the hyperparameters
- This means validation performance does not 100% truly reflect the performance for unseen data set (as our model was optimized for the validation performance)
- **Test data set** is different from the validation data set in the sense that it never anyhow affected the training procedure

Summary

- Model selection problem
- Bias-variance tradeoff
 - If we decrease one, the other will increase
- Underfitting & Overfitting
 - Model capacity, model complexity
- Generalization & Validation
- Splitting dataset
 - Training, validation, test

References

- Lecture notes
 - CC229 lecture note
 - <http://cs229.stanford.edu/notes2020fall/notes2020fall/cs229-notes5.pdf>
- Website
 - CS231n course website: <https://cs231n.github.io/>