

# **CoE202**

## **Fundamentals of Artificial intelligence**

### **<Big Data Analysis and Machine Learning>**

## **Deep Learning with libraries**

Prof. Young-Gyu Yoon  
School of EE, KAIST

# Contents

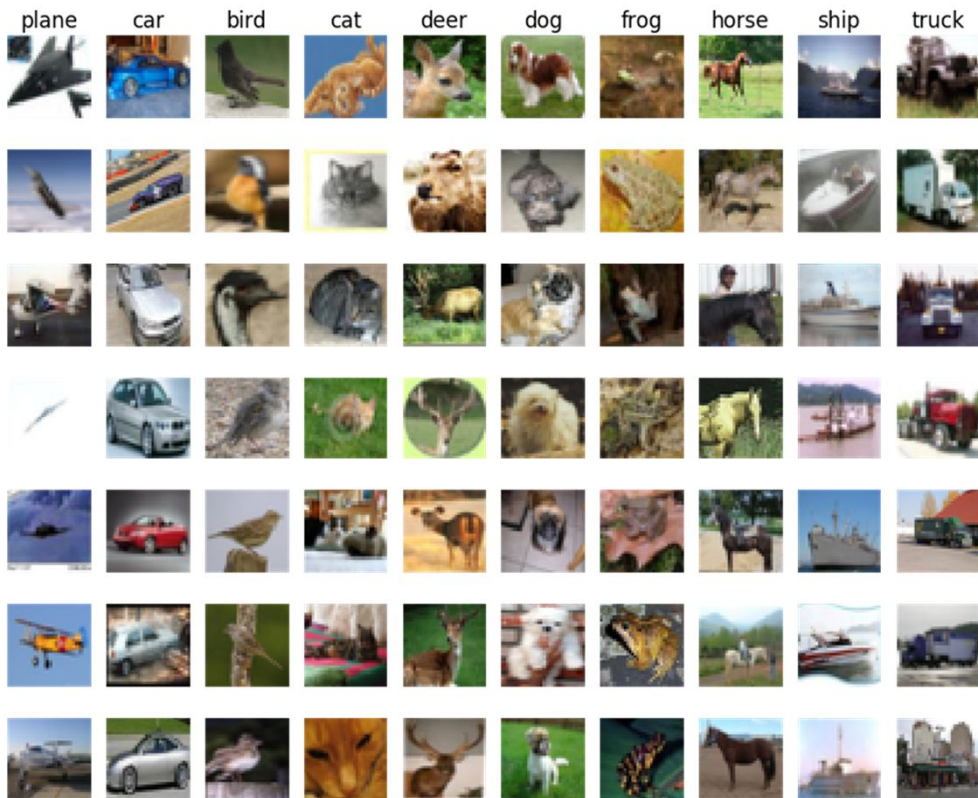
- Recap
  - Backpropagation
- Deep learning with libraries
  - Libraries
  - What we (human) do and what library does
  - Task formulation
  - NN design flow

# With deep learning



- MNIST digit recognition
  - 70,000 images
  - Image size: 28 x 28 pixels
  - Number of classes: 10

# With deep learning



- CIFAR-10 image classification
  - 60,000 images
  - Image size: 32 x 32 pixels
  - Number of classes: 10

# Deep learning libraries



- Developed by Google
- Good for professional use



- Developed by Facebook
- Easy to use



- Developed by University of Montreal
- Good for professional use



- Developed by a Google engineer
- Interface for Tensorflow, R, Theano, ...



- Developed by University of California, Berkeley
- Caffe2 was merged into PyTorch

# Why use library?

- **For example, with PyTorch**

- Complicated “forward function” can be easily implemented by **calling built-in functions**
- Lots of **loss functions** are implemented. We can just call them without having to worry about implementation
- Lots of **optimization methods** are included (more than just SGD!)
- **GPU acceleration** is supported
- Supports **Autograd**
  - We just need to define the “forward propagation”
  - We don’t have to worry about complicated gradient calculation anymore. **Backpropagation is automatically done** for us 😊

# Why use library?

- **For example, with PyTorch**

- We can **download lots of datasets** through torchvision that are suited for studying deep learning
- **Popular network architectures** are included
- Useful (and common) **image transformation** functions are included
- Come with data loading & shuffling functions
- ... and a lot more

# Why use library?

- **Simply put, Pytorch can do (almost) everything for us**
  - It allows us to implement neural networks without even understanding
    - How to do backpropagation and calculate gradient
    - How GD (and other optimization method) works
    - How to calculate a loss function
    - ...and the list goes on



# Why use library?

- Does it really mean we don't have to know anything?
- Then, why are we “wasting” our time learning all this?
- The library
  - **allows** us to focus on the concept and innovation (rather than details inside)
  - does **NOT** debug itself: we still have to understand how things work (in general) to be able to fix, debug and improve things

# PyTorch vs. Numpy

```
class mlp_classifier(torch.nn.Module):
    # initialization
    def __init__(self):
        super(mlp_classifier, self).__init__()
        self.layer1 = torch.nn.Linear(2, 10)
        self.layer2 = torch.nn.Linear(10, 10)
        self.layer3 = torch.nn.Linear(10, 10)
        self.layer4 = torch.nn.Linear(10, 1)
        self.relu = torch.nn.ReLU()
```

Define network

```
# forward path
def forward(self, x):
    # Used relu to add non-linearity
    x = self.layer1(x)
    x = self.relu(x)
    x = self.layer2(x)
    x = self.relu(x)
    x = self.layer3(x)
    x = self.relu(x)
    x = self.layer4(x)
    return x
```

Define forward propagation

With PyTorch

vs.

```
class nonlinear_classifier():
    # set initial weights
    def __init__(self, W1_init, W2_init):
        super(nonlinear_classifier, self).__init__()
        self.W1 = W1_init
        self.W2 = W2_init

    # forward pass
    def forward(self, x):
        # xk : input, zk: Wk*xb, yk = h(zk) (h is the activation function) @ kth layer
        self.x1 = x # input of the 1st layer = input of the network
        self.y1, self.z1 = self.forward_single_layer(self.W1, self.x1)
        self.x2 = self.y1 # input of the 2nd layer is the output of the first layer
        self.y2, self.z2 = self.forward_single_layer(self.W2, self.x2)
        y = self.y2 # output of the network = output of the 2nd layer

        return y # return network output

    def backward(self, label):
        y = self.y2

        # calculate loss and accuracy
        loss = cross_entropy(y, label)
        loss_avg = np.mean(loss)
        prediction_threshold = (y>0.5)
        accuracy = np.mean(prediction_threshold==label)
        n_sample = label.shape[1]

        # back-propagate
        dLdy = (y-label)/(y*(1-y))
        dLdz, dLdW2 = self.backward_single_layer(self.W2, self.x2, self.y2, dLdy)
        dLdy, dLdW1 = self.backward_single_layer(self.W1, self.x1, self.y1, dLdy)

        return loss_avg, accuracy, dLdW1, dLdW2

    def forward_single_layer(self, W, x, y, dLdy):
        # x : input, z: W*xb, y = h(z) (h is the activation function)
        x_pad = np.concatenate((x, np.ones((1, x.shape[1]))), axis=0) # x_pad = [x; 1]
        z = np.matmul(W, x_pad) # z = [W b] X x_pad
        y = 1/(1 + np.exp(-z)) # y = sigmoid(z)

        return y, z

    def backward_single_layer(self, W, x, y, dLdy):
        n_data = x.shape[1]

        dLdz = dLdy*y*(1-y) # backprop sigmoid
        x_pad = np.concatenate((x, np.ones((1, x.shape[1]))), axis=0) # add ones (for bias term)
        dzdx = x_pad.T # take transpose
        dzdW = W[:, :-1].T # exclude bias part then transpose

        dLdW = np.matmul(dLdz, dzdW)/n_data # divide by the number of data points (average)
        dLdx = np.matmul(dzdx, dLdz)

        return dLdx, dLdW

    # update parameters
    def update(self, dW1, dW2):
        self.W1 = self.W1 + dW1
        self.W2 = self.W2 + dW2
```

Initialize parameters

Define forward propagation

Define backpropagation


Define single layer forward propagation

Define single layer backpropagation

Parameter update  
(GD algorithm not included)

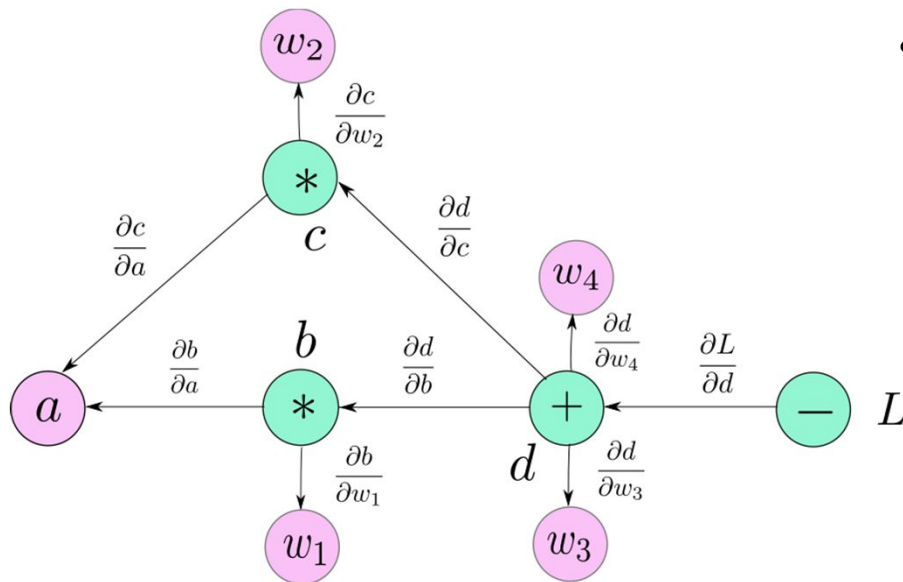
Using just Numpy

# PyTorch





- **What users do:**
  - Define
    - the forward function (equivalent to a network or a directed graph)

# PyTorch



- **What PyTorch does:**
  - Calculate all partial derivatives and do backpropagation
  - ...and update parameters using the algorithm and loss we choose
    - Of course we can build and use our own training algorithms and/or loss functions

# MINIST classification



$f$

$0 \sim 9$

$Y$

$X$

# MINIST classification

- We want to build a function (neural network) that maps ...

$$f(\text{8}) = 8$$

$$f(\text{0}) = 0$$

# MINIST classification



**Input:** 28 x 28 digital image (or 28 x 28 array)



**Output:** 10x1 vector (or 10 x 1 array), one-hot encoded

0	1	2	3		9
1	0	0	0		0
0	1	0	0		0
0	0	1	0		0
0	0	0	1		0
0	0	0	0	...	0
0	0	0	0		0
0	0	0	0		0
0	0	0	0		0
0	0	0	0		0
0	0	0	0		1

**Loss:** cross entropy loss (as we did for linear classification)

**Beware:** If you are using Pytorch and `torch.nn.CrossEntropyLoss` to calculate the loss, you should **NOT** convert the ground truth (target) label to one-hot encoded vector. It simply takes a integer number as the ground truth label.

# What about CIFAR-10 classification?



plane	car	bird	cat	...		truck
1	0	0	0			0
0	1	0	0			0
0	0	1	0			0
0	0	0	1			0
0	0	0	0			0
0	0	0	0			0
0	0	0	0			0
0	0	0	0			0
0	0	0	0			1

**Input:** 32 x 32 digital image (or 32 x 32 x 3 array)


**Output:** 10x1 vector (or 10 x 1 array), one-hot encoded

**Loss:** cross entropy loss (as we did for linear classification)


**Almost everything is the same as MNIST classification ...**



# Formulation: MNIST classification




# Formulation: MNIST classification



# Flatten?

1	1	0
4	2	1
0	2	1



1
1
0
4
2
1
0
2
1

## Flattening:

reshaping an array into a 1-dimensional array


$$3 \times 3 \rightarrow 9 \times 1$$

$$28 \times 28 \rightarrow 784 \times 1$$

$$15 \times 15 \times 3 \rightarrow 675 \times 1$$

# Formulation: MNIST classification


Neural Network



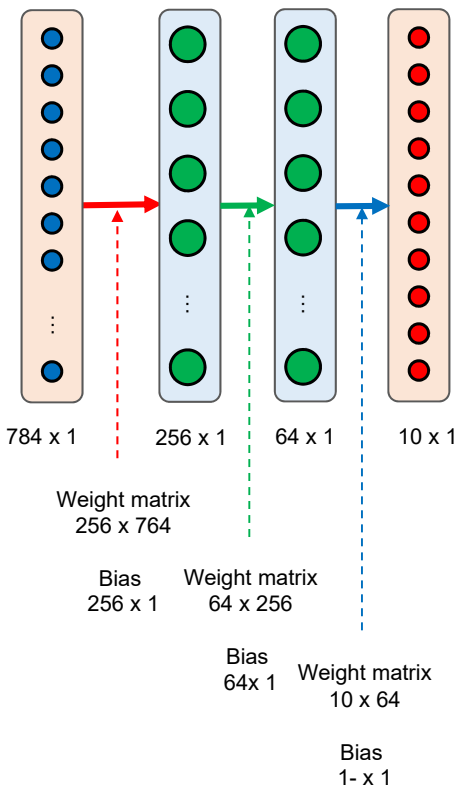
For MNIST classification,  
We need to design a network  
that takes 784x1 array as the input  
and generates 10x1 array as the output  
(**design constraint**)

There are infinitely many possible options

# Formulation: MNIST classification



# Neural Network Design



## Example design

Input: 784 x 1



1st layer: 256 x 1 neurons

2nd layer: 64 x 1 neurons

3rd layer (output): 10 x 1 neurons


How do we know if this is a good design?

# Neural Network Design: Bias-Variance Trade-off



**We want our network to have just right complexity!**

# Neural Network Design: Bias-Variance Trade-off




- This network has **3 layers**
- **Widths are 256, 64**
- **Number of parameters:**
  - Weight matrix
    - $256 \times 784 = 200,704$
    - $64 \times 256 = 16,384$
    - $10 \times 64 = 640$
  - Bias
    - 256, 64, 10
  - Total:  
 $200,704 + 16,384 + 640 + 256 + 64 + 10 = \mathbf{218,058}$

Ok, this seems to tell us something about the complexity ...




# Neural Network Design: depth vs. width

Training loss



iterations


Validation loss



iterations



- **Increasing depth** (number of layers) usually results in “faster” capacity increase than increasing width
- A network with a **smaller number of parameters** tends to be less susceptible to overfitting

# Neural Network Design: depth vs. width



- We want to use different form of parameterized function  $f(x; \theta)$  depending on our target function  $g(x)$  (which we do not know)
- Some families of parameterized functions  $f(x; \theta)$  can approximate  $g(x)$  with a smaller error with smaller number of parameters than others
- ...and this of course depends on  $g(x)$  (task-dependent)

# Neural Network Design: depth vs. width




**These are just conceptual plots!!!**

**Take-home message:**


The capacity (capability of approximate complex functions) and susceptibility to overfitting of different parameterized functions (e.g., neural network) scale differently (and this is task-dependent)

We are essentially trying to approximate a function. If we can reduce the degree of freedom of our parameterized function  $f(x; \theta)$  without compromising its capability to approximate the “target function,” then it would be desired

# Neural Network Design: avoid abrupt decrease




Size is gradually decreasing. Why?



Input data is being gradually “transformed” via multiple rounds of matrix multiplication & activation, so that it becomes **linearly separable** at the end

(why does it have to become linearly separable?)


# Neural Network Design: avoid abrupt decrease




- Neural network is not a magic
- Single layer can only do “ $h(WX)$ ”-much” of operation
- If we are asking a neural network to approximate a complicated function, we have to “give” enough depth to do that

Input data is being gradually “transformed”  
via multiple rounds of matrix multiplication & activation,  
so that it becomes **linearly separable** at the end


# Neural Network Design: avoid abrupt decrease



What is wrong with this design?




Is this OK?



Is this better?

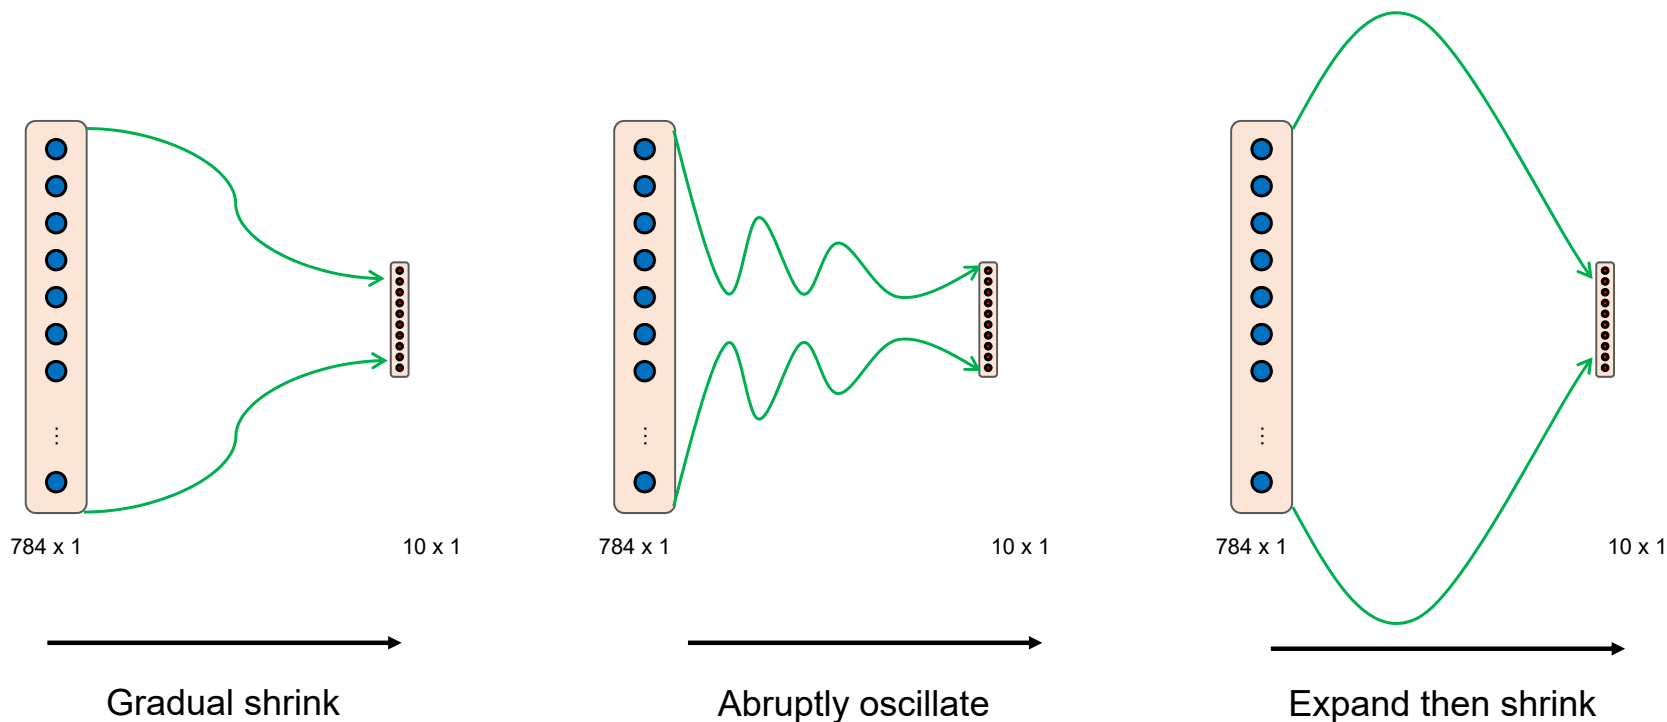
# Neural Network Design: avoid abrupt decrease



- Mapping a high-dimensional vector to a low-dimensional vector always comes with information loss, and the question is, whether the information that is relevant for estimating the output has been lost or not.
- Reducing dimension, without loss of (relevant) information, requires that the network is capable of “extracting” important information

From here to there is just a two-layer neural network which means it is unlikely that this two-layer network is capable of efficiently mapping the high-dimensional input to a low-dimensional form (efficiently: with a very small amount of information loss)




# Neural Network Design: avoid abrupt decrease





# How do we ever want to increase width?


- If the output has lower dimension than the input, why do we ever want to increase width?
  - Short answer: representing low dimensional data in a high dimension can be useful to implement certain operations (or functions)



# Neural Network Design: flow

- However, these design approaches still do not provide what actual depth and width (and other hyper-parameters) we should use
- **Possible Flow**
  - Build the first version then train (start with a small network)
  - Check training loss and validation loss
  - Increase the capacity until training loss sufficiently goes down
  - Stop increasing (or decrease) the capacity if the difference between training loss and validation loss becomes large
- The flow above is just one possible way (there are many other possible ways and many other things to consider)

# Neural Network Design: PyTorch Implementation



```
class mlp_classifier(nn.Module):  
    def __init__(self):  
        super(mlp_classifier, self).__init__()  
        self.layer1 = nn.Linear(28*28, 256)  
        self.layer2 = nn.Linear(256, 64)  
        self.layer3 = nn.Linear(64, 10)  
        self.relu = nn.ReLU()  
  
    def forward(self, x):  
        x = x.view(-1, 28*28)  
        x = self.layer1(x)  
        x = self.relu(x)  
        x = self.layer2(x)  
        x = self.relu(x)  
        x = self.layer3(x)  
        x = self.relu(x)  
        return x
```

# Summary

- Deep learning library makes it very easy to implement complicated neural networks
- ...but we still have to understand how they work inside to fully utilize the power of NN
- We have to formulate the given tasks in a form that NN can handle
- NN design involves lots of heuristics (trial and error), but there still are a few things to remember

# References

- Website
  - CS231n course website: <https://cs231n.github.io/>
- Pytorch tutorial
  - [https://pytorch.org/tutorials/beginner/deep\\_learning\\_60min\\_blitz.html](https://pytorch.org/tutorials/beginner/deep_learning_60min_blitz.html)
- Tensorflow guide
  - <https://www.tensorflow.org/guide/>
- Matplotlib Guide
  - <https://matplotlib.org/>
- Neural Networks, Manifolds, and Topology
  - <https://colah.github.io/posts/2014-03-NN-Manifolds-Topology/>