

# **CoE202**

## **Fundamentals of Artificial intelligence**

### **<Big Data Analysis and Machine Learning>**

## **Neural Networks Techniques**


Prof. Young-Gyu Yoon  
School of EE, KAIST

# Contents

- Recap
  - Biological vision
  - Convolution operation
  - ConvNet as a special case of NN
  - ConvNet as a generalization of NN
  - Building block of ConvNet
- Regularization methods
- Optimization methods
- NN architectures

# Regularization in Optimization

- **Regularization:** is the process of adding information in order to solve an ill-posed problem or to prevent overfitting




Green: with regularization

Blue: without regularization



- Modify loss function
- Give constraints to the network
- Decrease network size
- Quit training before the network over-fits
- Add “noise” to data

# Data augmentation




- Shift, crop, rotate, brightness adjustment (and more)
  - As long as we believe it is realistic enough ...

# Data augmentation?



- Remember that even human vision is orientation dependent
  - Rotation (and other operations) may or may not alter the content, so we need to think twice before we do it

# Data augmentation?



- The examples illustrate why we should think twice and check if the transformation may alter the contents!

# L2 (or L1) regularization

- **Problem**

- While training, the optimizer may tend to increase the absolute value of the parameters
- This may result in overfitting and/or floating point overflow

- **Solution:** modification of loss function

$$\mathcal{L}_{new} = \mathcal{L}_{original} + \lambda \mathcal{L}_{L2}$$


Add another term in our loss function to penalize for large absolute values of parameters

$$\mathcal{L}_{L2} = \sum_i w_i^2$$


$$-\nabla \mathcal{L}_{new} = -\nabla \mathcal{L}_{original} - \lambda \nabla \mathcal{L}_{L2}$$

$$-\nabla \mathcal{L}_{L2} = -2 \sum_i w_i$$

# Dropout



(a) Standard Neural Net




(b) After applying dropout.

- A simple technique to prevent overfitting
- (while training) we randomly “sub-sample” part of the neurons in network
- (while testing) we use all of the neurons



# Dropout



Q) Is this training accuracy or validation/test accuracy?

- **Training**


- We assign probability  $p$  to each neuron (the probability that each neuron will be present or not)
- Then, we train the network while randomly sub-sampling it
  - Randomly sub-sampling is based on the probability  $p$
  - We are training a network in a way such that 'a randomly selected part of the network' will be able to solve the task

- **Testing**


- For validation & testing, we enable all neurons

# Dropout

## 7.1 Effect on Features



(a) Without dropout



(b) Dropout with  $p = 0.5$ .

# Batch Normalization


**Input:** Values of  $x$  over a mini-batch:  $\mathcal{B} = \{x_1, \dots, x_m\}$ ;  
Parameters to be learned:  $\gamma, \beta$   
**Output:**  $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$
$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$
$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$
$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

**Algorithm 1:** Batch Normalizing Transform, applied to activation  $x$  over a mini-batch.

## Batch normalization

Ensure the output statistics of a layer are fixed.



# Batch Normalization

**Input:** Values of  $x$  over a mini-batch:  $\mathcal{B} = \{x_{1\dots m}\}$ ;  
Parameters to be learned:  $\gamma, \beta$   
**Output:**  $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$
$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$
$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$
$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

**Algorithm 1:** Batch Normalizing Transform, applied to activation  $x$  over a mini-batch.


- **Training**

- For each mini-batch, the mean and variance is calculated and the data is normalized
- The normalized data is scaled and shifted (linear transform) with trainable parameters  $\gamma$  and  $\beta$

- **Testing**

- For validation & testing, we use mean and variance from the entire dataset (not from a batch)

# Model Ensembles



- **Problem:**

- There always is a bias in our output (especially with high capacity models)

- **Solution:**

- Train multiple networks and take the averaged output (assuming that the bias will average out)

# RMSprop

- **Problem:**

$$\theta^{(k+1)} = \theta^{(k)} - \gamma \nabla \mathcal{L}(\theta^{(k)})$$

- the size of the gradient (used for gradient descent) may vary widely in magnitudes
  - some parameters experience large updates, some do not
  - proper learning rate is different from different parameters

- **Solution:** use adapter learning rates for different parameters

- divide the gradient (for each parameter) by the square root of the moving average of the squared gradient

$$MA^{(k)} = \beta MA^{(k-1)} + (1 - \beta)(\nabla \mathcal{L}(\theta^{(k)}))^2$$

“moving average” of squared gradient

$$\theta^{(k+1)} = \theta^{(k)} - \frac{\gamma}{\sqrt{MA^{(k)}}} \nabla \mathcal{L}(\theta^{(k)})$$

divide by the square root

# Adam: combine momentum and RMSprop

---

**Algorithm 1:** *Adam*, our proposed algorithm for stochastic optimization. See section 2 for details, and for a slightly more efficient (but less clear) order of computation.  $g_t^2$  indicates the elementwise square  $g_t \odot g_t$ . Good default settings for the tested machine learning problems are  $\alpha = 0.001$ ,  $\beta_1 = 0.9$ ,  $\beta_2 = 0.999$  and  $\epsilon = 10^{-8}$ . All operations on vectors are element-wise. With  $\beta_1^t$  and  $\beta_2^t$  we denote  $\beta_1$  and  $\beta_2$  to the power  $t$ .

---

**Require:**  $\alpha$ : Stepsize

**Require:**  $\beta_1, \beta_2 \in [0, 1)$ : Exponential decay rates for the moment estimates

**Require:**  $f(\theta)$ : Stochastic objective function with parameters  $\theta$

**Require:**  $\theta_0$ : Initial parameter vector

$m_0 \leftarrow 0$  (Initialize 1<sup>st</sup> moment vector)

$v_0 \leftarrow 0$  (Initialize 2<sup>nd</sup> moment vector)

$t \leftarrow 0$  (Initialize timestep)

**while**  $\theta_t$  not converged **do**

$t \leftarrow t + 1$

$g_t \leftarrow \nabla_{\theta} f_t(\theta_{t-1})$  (Get gradients w.r.t. stochastic objective at timestep  $t$ )

$m_t \leftarrow \beta_1 \cdot m_{t-1} + (1 - \beta_1) \cdot g_t$  (Update biased first moment estimate)

$v_t \leftarrow \beta_2 \cdot v_{t-1} + (1 - \beta_2) \cdot g_t^2$  (Update biased second raw moment estimate)

$\hat{m}_t \leftarrow m_t / (1 - \beta_1^t)$  (Compute bias-corrected first moment estimate)

$\hat{v}_t \leftarrow v_t / (1 - \beta_2^t)$  (Compute bias-corrected second raw moment estimate)



$\theta_t \leftarrow \theta_{t-1} - \alpha \cdot \hat{m}_t / (\sqrt{\hat{v}_t} + \epsilon)$  (Update parameters)

**end while**

**return**  $\theta_t$  (Resulting parameters)


---

# Comparison of optimization methods






# Comparison of optimization methods




# Skip Connection (ResNet)




- **Problem:** gradient vanishing
  - The deeper, more the powerful
  - However, a very deep network is difficult to train
  - Back-propagation through very large number of layers needs to be done  $\rightarrow$  it becomes difficult to properly train “front-end” layers
- **Solution:** make skip connections
  - Add direction connection paths in our network, so back-propagation can be done without gradient vanishing problem

# Skip Connection (ResNet)




**Without skip connections**

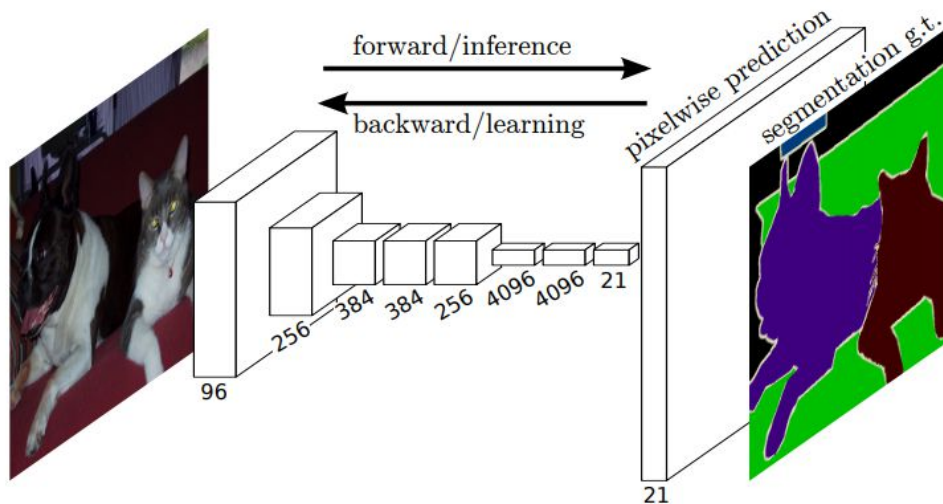


**With skip connections**

# AlexNet: winner of ImageNet challenge 2012

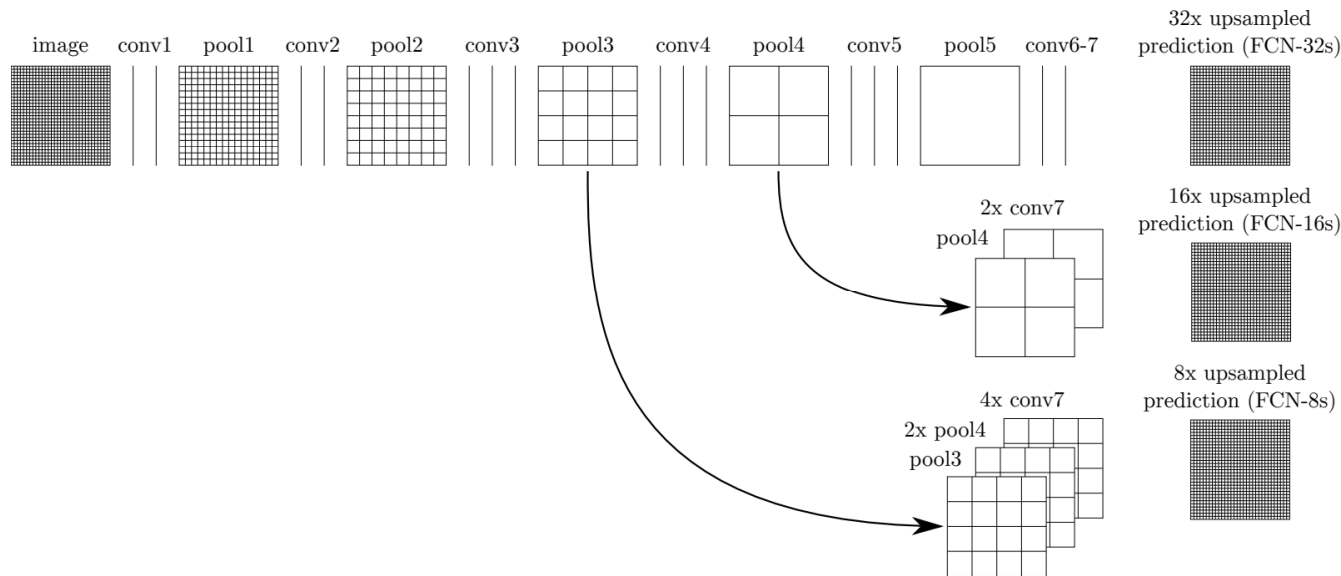


# Fully Convolutional Network




- We want more than just image classification  $\rightarrow$  we want to know what's where
- Conceptually speaking, we want  $X: 512 \times 512 \times 3 \rightarrow Y: 10 \times 512 \times 512$ , instead of  $X: 512 \times 512 \times 3 \rightarrow Y: 10 \times 1$  mapping

# Fully Convolutional Network




- We want more than just image classification → we want to know what's where
- Conceptually speaking, we want  $X: 512 \times 512 \times 3 \rightarrow Y: 10 \times 512 \times 512$ , instead of  $X: 512 \times 512 \times 3 \rightarrow Y: 10 \times 1$  mapping

# U-net



- Network architecture suited for  $X: 512 \times 512 \times 3 \rightarrow Y: 10 \times 512 \times 512$

# U-net



- **Key properties**
  - **Skip-connection:** high resolution information is directly fed-through
  - **Contraction-and-expansion:** abstraction (high level information extraction) is performed and used for pixel level segmentation
  - **Large receptive field:** convolution operations for image images corresponds to large-kernel convolution for full-sized images
  - **Multi-resolution processing:** information is processed at multiple resolutions



# Summary

- Regularization methods
  - Data augmentation
  - L1/L2 regularization
  - Dropout
  - Batch normalization
  - Model ensemble
- Optimization methods
  - RMSprop
  - Adam
- NN architectures
  - ImageNet
  - FCN
  - U-net

# References

- Website
  - CS231n course website: <https://cs231n.github.io/>